

# 12410106 Assignment1

Zhang Sihua

DeepLearning 25 fall

**Abstract.** This is the report about assignment1 of DeepLearning course, perceptron&MLP . In this report, we illustrate the implementation details and results analtsis of the assignment. The report contains three part: the perceptron, the mutli-layer perceptron and stochastic gradient descent. To run our code, just run the jupyter notebook in directory Part\_x (x=1,2,3) or execute bash commands provided in the notebook to run python scripts.

## 1 Part1 perceptron

In this part we implement and test a simple artificial neuron: a perceptron.

### 1.1 Task1

First we generate a dataset of points in  $R^2$  from two Gaussian distributions and sample 100 points from each. We keep 80 points per distribution as training data(160 in total), 20 for testing (40 in total).

We specify the mean and covariance between two features of each Gaussian distribution, and use `numpy.random.multivariate_normal()` to sample from each two-dimensional Gaussian distribution. Then we use `numpy.ones()`to generate labels. Two Distributions serve as two class, denoted as -1 and 1 respectively. Finally we stack them together as  $X$  and  $y$ , and visualize them using matplotlib. The visualization result is Fig.1.

Then, we use `train_test_split()` in sickit-learn to split the data into training and testing sets. The splitting ratio is 4:1 for training and testing.

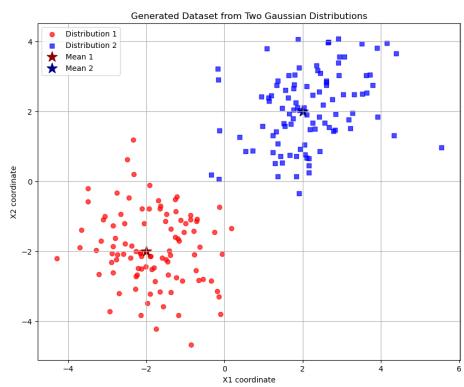
### 1.2 Task2

We implement the perceptron following the specs in `perceptron.py` and the pseudocode in `perceptronslides.pdf`. Note that I add a `batch_size` parameter to support mini-batch gradient descent because it is both efficient and robust, and a test function for testing convenience. Also I concatenate  $W, b$  together as shown in `perceptron_tutorial.pdf`. Implementation details can be found in source code.

### 1.3 Task3

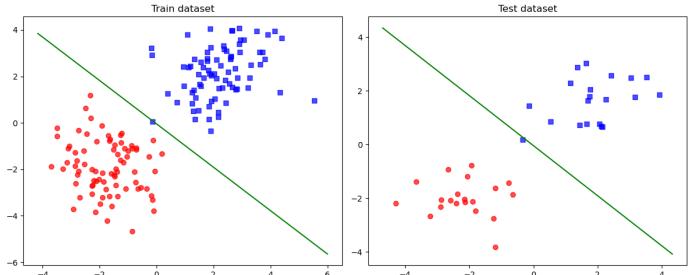
We train the perceptron on the training data (160 points) and test on the remaining 40 test points. The test accuracy is 0.97. We further visualize the decision boundary to inspect classification outcome. The visualization result is Fig.2.

045  
046  
047  
048  
049



**Fig. 1.** 2Gaussian

074  
075  
076  
077  
078  
079  
080  
081  
082  
083

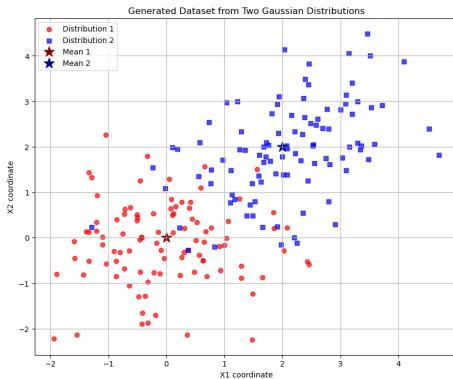


**Fig. 2.** DecisoinBoundary

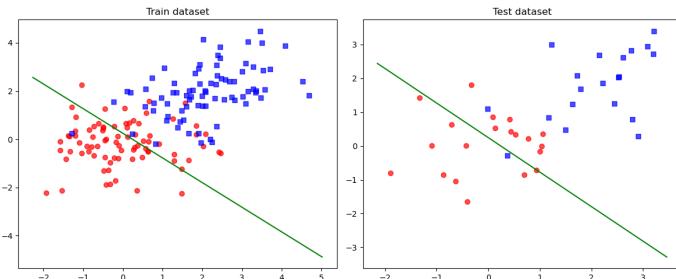
## 090 1.4 Task4

091 We do more experiments on different means and variances of gaussian distributions and analyse the result. Intuitively, if the two Gaussians' means are two  
 092 close or their variance is too large, they may not be linearly separable since the  
 093 overlapping of probability density become larger.  
 094

095 Fig.3 and Fig.4 shows the result where the means of two distribution are  
 096 closer to each other, variances remain the same. Fig.5 and Fig.6 shows the result  
 097 where two variances both increase from 1 to 3, and means remains the same. We  
 098 can see in both cases, the decision boundary fails to perfectly classify all data  
 099 points correctly. This is because perceptron is a linear model, when the dataset  
 100 is no longer linearly separable, it is not able to find a separating hyperplane  
 101 for all data points. Divergence of loss during training is also a strong indicator  
 102 to this.



117 118 **Fig. 3.** CloserMeanDistributions



132 133 **Fig. 4.** CloserMeanDicision

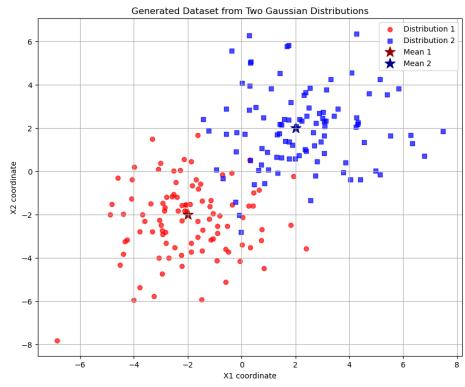
135  
136  
137  
138  
139

Fig. 5. LargerVarDistributions

150

151

152

153

154

155

156

157

158

159

160

161

162

163

164

165

166

167

168

169

170

171

172

173

174

175

176

177

178

179

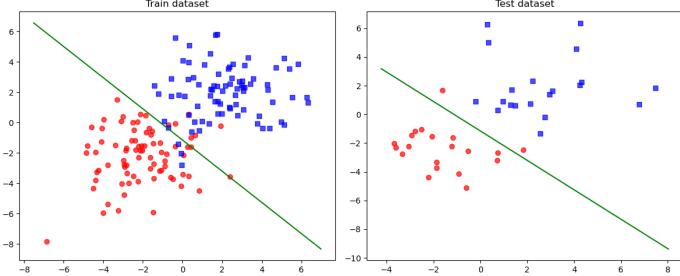


Fig. 6. LargervarDicision

174

175

176

177

178

179

135  
136  
137  
138  
139

140

141

142

143

144

145

146

147

148

149

150

151

152

153

154

155

156

157

158

159

160

161

162

163

164

165

166

167

168

169

170

171

172

173

174

175

176

177

178

179

## 180 2 Part2 multi-layer perceptron

### 181 2.1 Task1

182 We implement Linear, ReLU, Softmax and CrossEntropy in modules.py using  
 183 numpy. Keypoints are:

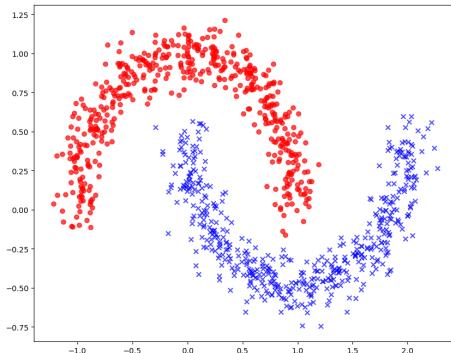
- 184 1. We should initialize a cache to store intermediate values. When we need  
 them to calculate gradients, we can directly fetch them from cache to avoid  
 recomputation.
- 185 2. Be careful with matrix calculus where errors can be easily introduced. For  
 example, the derivative of  $XW + b$  w.r.t  $W$  is  $X^T$ , not  $X$ .

186 Then we implement a simple MLP using these modules. Although our MLP  
 187 isn't going to be extremely deep, I still use He initialization for the weights in  
 188 Linear layers, which can prevent gradient from exploding or vanishing to some  
 189 extent. It is known that He initialization does better than Xavier initialization for  
 190 ReLU activation function. However, we still encounter the problem of gradient  
 191 explosion, probably due to the lack of normalization or large learning rate (in  
 192 our experiment is 0.01). But it does not degrade the whole model.

193 Besides, I add update\_params() function to update the parameters of the  
 194 network during training.

### 201 2.2 Task2

202 After our MLP is ready to use, we use scikit-learn's  
 203 `make_moons()`, `train_test_split()`, `OneHotEncoder` to generate and preproc-  
 204 cess our datasets. The visualization result is Fig.7.



219 **Fig. 7.** Moons

220 Then we implement training and testing script in train\_mlp\_numpy.py. I  
 221 prefer mini-batch gradient descent (MBGD) and didn't get the meaning that we

ought to implement batch gradient descent (BGD) in this part, hence I add a dataloader class to assist mini-batch data loading. However, you can find BGD, MBGD, SGD in Part3.

### 2.3 Task3

We test the MLP on test dataset with default parameters. If you unset the random seed and test it multiple times you should find that on average the accuracy is about 0.5, which is not better than randomly guessing since we only have 2 classes. Then we train it on train dataset using BGD for 15 epochs. Batch size is set to 32. We record the running losses and accuracies during training, and test it every 1 epoch (this hyperparameter can be larger than 1). Finally we visualize the loss and accuracy curve for trian and test dataset in Fig.8.

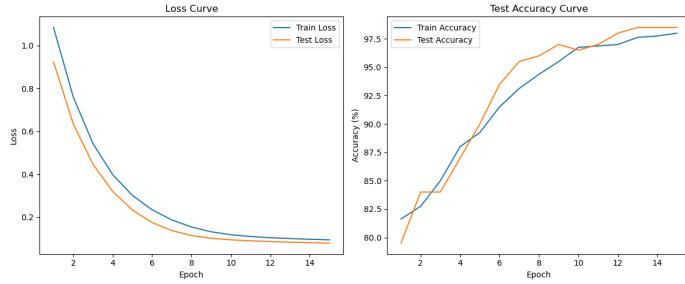


Fig. 8. LossAccCurves

The figure shows that at the beginning of training, train\_loss is high and train\_acc is low (80% this time, thanks to good random seed) . As the training progresses, train\_loss decreases and train\_acc increases. Our MLP gradually learns to classify the training data.

Note that sometimes the test\_loss is even lower than the train\_loss. This can happen due to random data splits, or the test set being easier than the training set by chance. As training progresses, train\_loss usually becomes lower than test\_loss. This does not indicate overfitting or underfitting, but rather reflects the randomness in data and training process.

## 3 Part3 stochastic gradient descent

### 3.1 Task1

I misunderstood the instruction and implemented minibatch gradient descent in Part2. As a result, I implement batch gradient descent, minibatch gradient descent and stochastic gradient descent all three kinds of gradient descent algorithms in Part3.

The only changes from Part2 is that I add two new command line argument to specify the gradient descent mode, `-gd_mode`, along with some changes in the train function, and `-batch_size`, which is the batch size for mini-batch gradient descent.

`-gd_mode` can take three int values: 0 stands for "batch", 1 stands for "mini-batch", and 2 stands for "stochastic". The default value is 0.

`-batch_size` only work if `-gd_mode` is 1. The default value is 32.

Model and training configuration is same as Part2.

### 3.2 Task2

Again we test our MLP with default parameters to show that our model is initialized randomly. Then we use three gradient descent algorithms to train our model and plot losses and accuracies results. The results of BGD, MBGD, SDG is Fig.9, Fig.10, Fig.11. The training hyperparameters are

1. BGD: epochs = 200, learning\_rate = 1e-2
2. MBGD: epochs = 15, Batch\_size = 32, learning\_rate = 1e-2
3. SGD: epochs = 15, Batch\_size = 1, learning\_rate = 1e-3

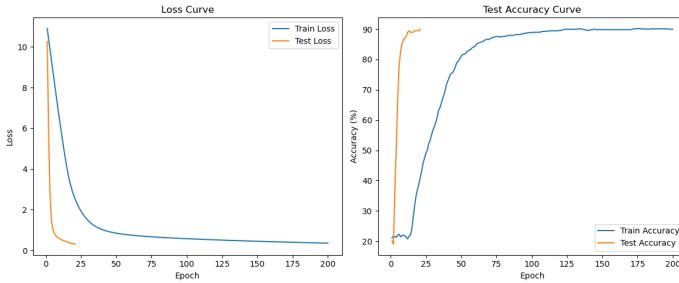


Fig. 9. BGD

Batch Gradient Descent (BGD) use all training examples as a batch, while mini-batch gradient descent (MBGD) uses a subset of training examples as a batch and Stochastic Gradient Descent (SGD) uses only one training example as a batch. When we decrease the batch size, our experiments show that the model will converge faster, but probably more unstable. If you run the notebook multiple times, you will see that BGD leads to similar loss and acc curves, while the curves of MBGD and SGD can be vary dramatically each time. This is because BGD use the whole dataset to learn, but mini-batch gradient descent randomly samples a subset to approximate the whole data distribution, and SGD iterates over a single example which add even more noise. Adding noise to training process may result in difficulty in converging, but it can also help to avoid overfitting and improve the model's generalization ability.

315  
316  
317  
318  
319  
320

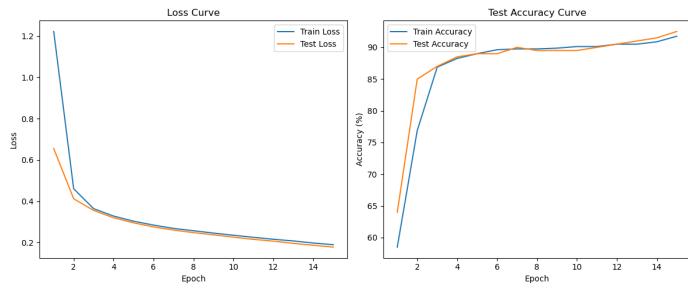


Fig. 10. MBGD

330  
331  
332  
333  
334  
335  
336  
337  
338  
339  
340  
341  
342

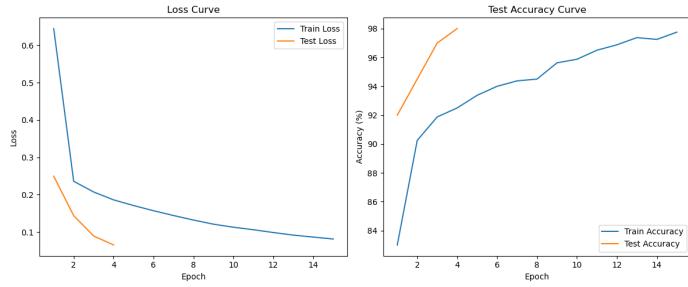


Fig. 11. SGD

353  
354  
355  
356  
357  
358  
359

315  
316  
317  
318  
319  
320  
321  
322  
323  
324  
325  
326  
327  
328  
329  
330  
331  
332  
333  
334  
335  
336  
337  
338  
339  
340  
341  
342  
343  
344  
345  
346  
347  
348  
349  
350  
351  
352  
353  
354  
355  
356  
357  
358  
359